

# NLP with Python

Using NLTK and SpaCy for text classification

## AGENDA

1. Introduction to Text Classification
2. Overview of Nltk and SpaCy
3. Text classification in Nltk
4. Text classification in spaCy

**“You were born with wings, why prefer to crawl through life?”**  
**- Rumi**

What does this mean to you? How might AI interpret ‘wings’ literally and miss the deeper meaning?

## 1. Introduction to Text Classification

Text classification is a Natural Language Processing (NLP) task where text is analyzed, understood, and categorized into predefined classes or categories.

It involves assigning a label or category to a given piece of text based on its content.

- Categorizing text into predefined categories based on its content.
- **Examples:**
  - Email filtering (spam vs. not spam).
  - Sentiment analysis (positive, negative, neutral).
  - News categorization (politics, sports, technology).

- **Importance:** Automates the understanding of large volumes of text.

## Applications of Text Classification

- **Sentiment Analysis:** Classifying text as positive, negative, or neutral (e.g., analyzing customer reviews).

Spam Detection: Identifying spam or non-spam emails.

- **Topic Categorization:** Grouping articles or documents by topic (e.g., politics, sports, technology).
- **Language Detection:** Identifying the language of a given text.
- **Intent Detection:** Understanding the intent behind a query in chatbots or virtual assistants.

## Methods Used in Text Classification

- **Rule-Based Systems:** Predefined rules are used to classify text based on patterns or keywords. For example, an email containing the phrase "Congratulations, you've won!" could be classified as spam.
- **Machine Learning Models:** Algorithms such as *Naive Bayes*, *Support Vector Machines (SVM)*, or *Logistic Regression* are trained on numerical features extracted from text data. These models learn patterns from labeled datasets to make predictions.
- **Deep Learning Models:** Advanced neural network architectures like Convolutional Neural Networks (CNNs), or Transformers (e.g., BERT, GPT) are used for understanding and classifying complex text data, especially when large datasets are available.

## 2. Overview of NLTK and SpaCy

### NLTK

- **NLTK:**
  - Python library for working with text.
  - Designed primarily as an educational tool for teaching and learning NLP.
  - For basic NLP tasks like tokenization, stopword removal, stemming, lemmatization, POS tagging, and syntactic parsing (where a sentence is analyzed to determine its grammatical structure).
  - Generally slower because it processes text one step at a time and is designed more for research and teaching.
- NLTK can be slower when dealing with large datasets or real-time processing due to its older design.

- [NLTK Website](#)
- [NLTK Book](#)
- [NLTK Sentiment Analysis](#)

Building a Text Classification Pipeline with NLTK

1. **Import libraries.**
2. **Load and preprocess text data.**
3. **Feature extraction using Bag of Words or TF-IDF.**
4. **Train a simple Naive Bayes classifier.**
5. **Test and evaluate the classifier.**

## spaCy

- **SpaCy:**
  - Built for production-level tasks and focuses on efficiency and performance. It is optimized for industrial use cases, handling large-scale data and real-time NLP applications.
  - Known for its speed and efficiency.
  - For advanced preprocessing like lemmatization, POS tagging, dependency parsing, named entity recognition (NER), and word embeddings.
- [spaCy 101](#) A brief introduction
- [spaCy Course](#) Advanced NLP with spaCy

Text Classification with SpaCy

1. **Load SpaCy and process text.**
2. **Extract linguistic features (POS, entities, etc.).**
3. **Train a classifier using features.**
4. **Evaluate the classifier.**

## 3. Text Classification with NLTK

```
In [ ]: import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.classify import NaiveBayesClassifier

# Download NLTK resources
# Download necessary datasets
nltk.download('movie_reviews') # Movie review dataset
nltk.download('stopwords')    # Stopwords for preprocessing
nltk.download('punkt')        # Tokenizer for text processing
```

We create a small dataset with sentences labeled as "pos" (positive) or "neg" (negative).

```
In [ ]: # Step 1: Sample dataset of sentences labeled as positive or negative
training_data = [
    ("I love this movie", "pos"),
    ("This film is amazing", "pos"),
    ("I hated this movie", "neg"),
    ("This film is terrible", "neg")
]
```

### Preprocessing:

- Tokenization: Break the sentence into individual words.
- Stopword Removal: Remove common words like "the", "and", etc., to focus on meaningful words.
- Feature Extraction: We create a dictionary where the keys are words, and the values are True (indicating the presence of the word in the sentence).

```
In [ ]: # Step 2: Preprocessing and Feature Extraction
def extract_features(sentence):
    words = word_tokenize(sentence.lower()) # Tokenize and convert to lowercase
    stop_words = set(stopwords.words('english')) # Get stop words
    words = [word for word in words if word.isalpha() and word not in stop_words]
    return {word: True for word in words} # Create feature dictionary with words

# Convert training data into feature sets
training_features = [(extract_features(sentence), label) for sentence, label in
```

Naive Bayes Classifier: We train the classifier on the feature set.

```
In [ ]: # Step 3: Train Naive Bayes Classifier
classifier = NaiveBayesClassifier.train(training_features)
```

Testing: We classify new sentences using the trained classifier.

```
In [ ]: # Step 4: Test the classifier with new sentences
test_sentences = [
    "I really enjoyed this movie",
    "This movie was awful"
]
```

```
In [ ]: for sentence in test_sentences:
    features = extract_features(sentence)
    predicted_label = classifier.classify(features)
    print(f"Sentence: '{sentence}' => Predicted Sentiment: {predicted_label}")
```

## 4. Text Classification with SpaCy

```
In [ ]: #pip install spacy
```

```
In [ ]: #pip install blis
```

```
In [ ]: !pip install spacy
```

```
In [ ]: import spacy
from spacy.pipeline.textcat import Config, ConfigSchema, TextCategorizer
from spacy.training.example import Example
```

```
In [ ]: #pip show spacy pydantic
```

```
In [ ]: #pip install "pydantic<2.0"
```

```
In [ ]: !pip install spacy
```

```
In [ ]: pip install --upgrade pip
```

```
In [ ]: pip show spacy
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]: # Load SpaCy's pre-trained English model
nlp = spacy.load('en_core_web_sm')

# Step 1: Create the text classifier component
text_cat = nlp.create_pipe('textcat', config={"exclusive_classes": True, "archit
nlp.add_pipe(text_cat, last=True)
```

```
In [ ]: # Add Labels (positive and negative)
text_cat.add_label("pos")
text_cat.add_label("neg")
```

```
In [ ]: # Step 2: Prepare the training data
# The data format is [(text, label), where label is 'pos' or 'neg']
training_data = [
    ("I love this movie", {"cats": {"pos": 1, "neg": 0}}),
    ("This film is amazing", {"cats": {"pos": 1, "neg": 0}}),
    ("I hated this movie", {"cats": {"pos": 0, "neg": 1}}),
    ("This film is terrible", {"cats": {"pos": 0, "neg": 1}})
]
```

```
In [ ]: # Convert training data to SpaCy's format
train_examples = []
for text, annot in training_data:
    doc = nlp.make_doc(text)
    example = Example.from_dict(doc, annot)
    train_examples.append(example)
```

```
In [ ]: # Step 3: Train the classifier
optimizer = nlp.begin_training()
for epoch in range(10): # 10 epochs for training
    losses = {}
    # Shuffle and batch the examples
    # Example usage of batching (make sure you shuffle the training data)
    # this is a simple approach with no batch strategy.
    for batch in spacy.util.minibatch(train_examples, size=2):
```

```
nlp.update(batch, losses=losses)
print(f"Epoch {epoch} Losses: {losses}")
```

In [ ]:

```
# Step 4: Test the classifier on new sentences
test_sentences = [
    "I really enjoyed this movie",
    "This movie was awful"
]

for sentence in test_sentences:
    doc = nlp(sentence)
    print(f"Sentence: '{sentence}' => Predicted Sentiment: {'positive' if doc.ca
```

In [ ]:

In [ ]:

## Comprehensive Text Classification with NLTK including feature engineering (Optional)

### What We Will Cover

1. Understanding Text Classification
2. Dataset Preparation
3. Preprocessing Steps
4. Feature Extraction
5. Training a Naive Bayes Classifier
6. Testing and Evaluating the Model

### 1. Understanding Text Classification

Text classification involves assigning a category to a given piece of text. For example:

- Labeling emails as "**Spam**" or "**Not Spam.**"
- Categorizing movie reviews as "**Positive**" or "**Negative.**"

In this exercise, we will classify movie reviews as either **Positive** or **Negative** using NLTK.

```
In [ ]: # Import necessary libraries
import nltk
from nltk.corpus import movie_reviews
from nltk.corpus import stopwords
from nltk.classify import NaiveBayesClassifier
from nltk.classify.util import accuracy
```

### 2. Dataset Preparation

NLTK comes with a built-in dataset for movie reviews called `movie_reviews`. Each review in this dataset is pre-labeled as positive or negative.

```
In [ ]: # Download necessary datasets
nltk.download('movie_reviews') # Movie review dataset
```

```

nltk.download('stopwords')      # Stopwords for preprocessing
nltk.download('punkt')          # Tokenizer for text processing

```

```

In [ ]: # Load Dataset
# Extract movie reviews and their associated labels
documents = [(list(movie_reviews.words(fileid)), category)
              for category in movie_reviews.categories()
              for fileid in movie_reviews.fileids(category)]

```

### 3. Preprocessing Steps

Before training the classifier, we need to preprocess the text:

1. **Tokenization:** Breaking text into individual words.
2. **Removing Stop Words:** Eliminating common, unimportant words like "the" and "is."
3. **Lowercasing:** Converting all words to lowercase to maintain consistency.

```

In [ ]: # Tokenization
# Example of tokenizing one document for clarity
sample_document = documents[0][0] # Get the first document (words)
print(f"Original Document: {sample_document[:20]}") # Print the first 20 words

# Lowercasing
# Convert all words to lowercase
lowercased_words = [word.lower() for word in sample_document]
print(f"Lowercased Words: {lowercased_words[:20]}")

# Removing Non-Alphabetic Tokens
# Remove words that are not purely alphabetic
alphabetic_words = [word for word in lowercased_words if word.isalpha()]
print(f"Alphabetic Words: {alphabetic_words[:20]}")

# Stop Word Removal
# Load stopwords and remove them from the dataset
stop_words = set(stopwords.words('english'))
filtered_words = [word for word in alphabetic_words if word not in stop_words]
print(f"Filtered Words (No Stopwords): {filtered_words[:20]}")

```

### 4. Feature Extraction

We will use the following method for feature extraction:

- Represent each review as a list of words and create a feature set where each word's presence is marked as **True** or **False**.
- The goal is to extract features from each document in the dataset after preprocessing. These features are required for training the Naive Bayes classifier.

```

In [ ]: #Extract Features

# Create a list of all words in the dataset
all_words = nltk.FreqDist(word.lower() for word in movie_reviews.words() if word)

# Use the 2000 most common words as features
word_features = list(all_words.keys())[:2000]

```

```

In [ ]: # Define a feature extractor function
def document_features(document):
    document_words = set(document)

```

```
features = {word: (word in document_words) for word in word_features}  
return features
```

```
In [ ]: # Preprocess all documents  
preprocessed_documents = []  
for (doc, category) in documents:  
    # Step 1: Lowercase and remove stopwords + non-alphabetic tokens  
    filtered_words = [word.lower() for word in doc if word.isalpha() and word.lc  
    # Step 2: Apply feature extraction  
    features = document_features(filtered_words)  
    preprocessed_documents.append((features, category))
```

## 5. Training a Naive Bayes Classifier

Naive Bayes is a simple yet powerful algorithm for text classification. It works well with small datasets and uses probabilities to predict the most likely category.

```
In [ ]: # Split the dataset into training and testing sets (80% training, 20% testing)  
train_set = preprocessed_documents[:1600]  
test_set = preprocessed_documents[1600:]
```

```
In [ ]: # Train the Classifier  
classifier = NaiveBayesClassifier.train(train_set)
```

## 6. Testing and Evaluating the Model

After training, we'll test the classifier on new data to measure its accuracy.

```
In [ ]: # Evaluate the Model  
print(f"Accuracy: {accuracy(classifier, test_set) * 100:.2f}%")
```

```
In [ ]: # Test with New Data  
test_review = "This movie was absolutely amazing, with great performances and a  
test_tokens = nltk.word_tokenize(test_review)  
test_words = [word.lower() for word in test_tokens if word.isalpha() and word.lc  
test_features = document_features(test_words)  
print(f"Prediction for test review: {classifier.classify(test_features)}")
```

```
In [ ]: # Display the Most Informative Features  
print("\nMost Informative Features:")  
classifier.show_most_informative_features(10)
```

Feature (e.g., "chick"): These are the words found in the text that are most useful for predicting sentiment.

Association (e.g., "neg : pos = 8.6 : 1.0"): This shows how strongly the word is linked to one class. For example, "chick" is 8.6 times more likely to appear in negative reviews than positive ones.

## Additional Tools used in NLP

### 1. Preprocessing Tools



- **NLTK (Natural Language Toolkit):** A comprehensive library for basic NLP tasks such as tokenization, stemming, lemmatization, and syntactic parsing.
- **SpaCy:** An efficient library for advanced NLP tasks like dependency parsing, named entity recognition (NER), and word embeddings.

## 2. Text Representation Tools

- **Gensim:** A library for creating word embeddings and topic modeling using techniques like Word2Vec and LDA.
- **Transformers (by Hugging Face):** Provides access to pre-trained models like BERT and GPT for text embeddings and advanced NLP tasks.

## 3. Data Manipulation and Visualization Tools

- **Pandas:** A powerful library for cleaning, organizing, and manipulating textual datasets.
- **Matplotlib/Seaborn:** Libraries for visualizing data trends, such as word frequencies and sentiment scores.
- **WordCloud:** A simple tool for creating visual word clouds to represent text data.

## 4. Machine Learning and Deep Learning Tools

- **Scikit-Learn:** A versatile library for training machine learning models like Naive Bayes and SVM on text data.
- **TensorFlow/Keras:** A deep learning framework for building and training neural networks for NLP tasks.
- **PyTorch:** A flexible deep learning library for implementing advanced NLP models.

## 5. Data Acquisition Tools

- **BeautifulSoup:** A library for web scraping and extracting text data from HTML pages.
- **Scrapy:** An advanced tool for building web scrapers to gather large amounts of textual data.
- **Tweepy:** A Python library for accessing Twitter data via the Twitter API.

## 6. Annotation Tools

- **Label Studio:** A user-friendly tool for manually annotating text datasets for training machine learning models.
- **Prodigy:** A more advanced annotation tool that incorporates active learning to improve dataset quality.

## 7. Cloud NLP Tools (Optional)

- **Google Cloud Natural Language API:** A cloud-based service for sentiment analysis, entity extraction, and more.
- **AWS Comprehend:** Amazon's NLP service for text classification, entity recognition, and sentiment analysis.

- **Azure Text Analytics:** A cloud tool for performing key phrase extraction, sentiment analysis, and more.

## Live Exercise

Now it's your turn!

### Task 1: description of task

- instructions

\_\_\_\_\_END\_\_\_\_\_



THANK YOU!

Live Exercise Solutions

In [ ]: solutions

## Programming Interview Questions

1. topic:

- question

In [ ]:

**Mohammad Idrees Bhat**

Tech Skills Trainer | AI/ML Consultant